

Performance Analysis and Optimization of C++ Standard Libraries

Sebastian Pop, Aditya Kumar

SARC: Samsung Austin R&D Center

October 26, 2017

Agenda

- ▶ C++ standard template libraries
- ▶ software performance analysis
- ▶ improvements to libc++ and libstdc++ performance

C++ Standard Template Libraries

- ▶ STL is easy to use
 - ▶ standard interface: portable
 - ▶ easy to change data types: list, vector, deque, map, etc.
 - ▶ easy to change algorithms: iterators
- ▶ complexity specified by standard
- ▶ performance left to implementation

Performance of STL implementations

- ▶ performance and memory usage depend on
 - ▶ implementation: libc++ vs. libstdc++ vs. MSVC, etc.
 - ▶ container type
 - ▶ inefficiencies in implementations
- ▶ always analyze software performance to validate your choice

Figure: Size in bytes of empty containers on x86_64

Container	libstdc++	libc++	MSVC
vector<int>	24	24	24
list<int>	24	24	16
deque<int>	80	48	40
set<int>	48	24	16
unordered_set<int>	56	40	64
map<int, int>	48	24	16
unordered_map<int, int>	56	40	64

Software Performance Analysis

- ▶ identify hot functions from execution profiles
- ▶ inspect hot path: unit-benchmarking
- ▶ identify resource utilization on hot path

Profiling: identify hot path

- ▶ linux-perf: cycles, instructions, HW counters
- ▶ valgrind: cachegrind (R/W/Instrs), callgrind
- ▶ oprofile

Unit-benchmarking: inspect hot path

unit-benchmarking is unit-testing for performance

- ▶ set-up data structures in memory
- ▶ call hot function
- ▶ execute hot function until performance measures stabilize ¹

check performance of a single hot operation: less noise, keep focus

¹<https://github.com/google/benchmark>

Analyze resource utilization on hot path

Inspect:

- ▶ source code, compiler IR, assembly code
- ▶ CPU usage, instructions used and their latencies
- ▶ memory bus and caches: loads/stores, spills, cache misses

Improve Software Performance

- ▶ eliminate unnecessary work
 - ▶ call functionality from libc or libc++
 - ▶ reduce bus traffic: vectorize loads and stores
 - ▶ help compiler remove redundancies: attributes and inline
- ▶ analyze performance of different implementations
 - ▶ change data structures
 - ▶ change algorithms
 - ▶ change STL implementations
- ▶ analyze trade-offs of caching previous results
 - ▶ use more memory vs. less computation (and vice versa)

Our contributions to improve performance of libc++ and libstdc++

- ▶ string to int value parsing: `xsgetn` in libc++
 - ▶ replace byte by byte copy with call to libc `memcpy`
 - ▶ important speedup on proprietary benchmark
- ▶ `std::string.find()` in libc++ and libstdc++
 - ▶ replace byte by byte compare with call to `memchr` and `memcmp`
 - ▶ 12x speedup on std-benchmark ²
- ▶ inline ctor/dtor
 - ▶ `shared_ptr`
 - ▶ `basic_string`
- ▶ add attribute `noreturn` to non-returning functions
 - ▶ `__locale`, `vector`, `deque`, `future`, `regex`, `system_error`, etc.
 - ▶ important for compiler optimizations
 - ▶ remove false positives in static analysis tools


²<https://github.com/hiraditya/std-benchmark>

std::sort (libc++)


- ▶ Convert to introsort
- ▶ Sorting technique, which begins with quicksort and switches to heapsort after recursion reaches a threshold
- ▶ Worst case complexity of $O(N\log N)$
- ▶ Eliminate recursion
- ▶ Replaced memory intensive recursive calls with stack
std::stack uses std::deque, which uses std::algorithm
- ▶ Improved worst case time complexity by a factor of 10
<https://reviews.llvm.org/D36423>
- ▶ quicksort with tail recursion elimination: quadratic worst case
- ▶ reimplemented as introsort: begin with quicksort, switch to heapsort when recursion depth goes beyond a threshold
- ▶ 16x speedup in the worst case (std-benchmark ³)

³<https://github.com/hiraditya/std-benchmark>

Issues with std::sort (libc++)



Sorting Results Plot (With std-benchmark)



Lessons learned (containers) optimizing destructor of string

```
#include<string>
```

```
int main() {  
    std::string s("a");  
    s+='a';  
    return 0;  
}
```

```
$ g++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv
```

```
$ clang++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv  
call    _ZdlPv
```

```
#include<string>  
void foo();
```

```
$ g++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv  
call    _ZdlPv
```

```
int main() {  
    const std::string s("a");  
    foo();  
    return 0;  
}
```

```
$ clang++ -O3 t.cpp -S -fno-exceptions -std=c++11 -o - | grep _ZdlPv
```

Annex

Performance Analysis with Valgrind

- ▶ `valgrind [--tool=memcheck]`
- ▶ `valgrind --tool=cachegrind`
 - ▶ cache and branch simulator
 - ▶ count read, write, and branch instructions
- ▶ `valgrind --tool=callgrind`
 - ▶ execution call graph
 - ▶ visualization tool `kcachegrind`
- ▶ `valgrind --tool=massif`
 - ▶ how much heap and stack memory your program uses
 - ▶ which parts of program allocate the heap memory

Valgrind: Example – SQLite

```
$ valgrind --tool=cachegrind ./sqlite_llvm <test.sql >/dev/null  
[...]
```

	Ir	IImr	ILmr	Dr	DImr	DLmr	Dw	DImw	DLmw	
1,278,771,731	29,231,219	35,783	359,414,267	6,707,514	528,920	197,515,528	2,594,262	171,968	PROGRAM TOTALS	

	Ir	IImr	ILmr	Dr	DImr	DLmr	Dw	DImw	DLmw	file:function
363,052,233	7,560,087	3,122	97,707,865	1,084,529	77,197	44,505,055	217,826	29,838	src/sqlite3.c:sqlite3VdbeExec	
95,048,357	80,721	111	33,248,107	59,086	7,273	20,173,275	91	7	src/sqlite3.c:vdbeRecordCompareWithSkip	
68,045,026	695,509	1,144	14,883,933	114,698	1,918	5,525,733	272,507	19,249	src/sqlite3.c:balance	
56,713,554	1,101,002	276	18,416,705	683,914	21,085	3,453,665	1,947	25	src/sqlite3.c:sqlite3BtreeMovetoUnpacked	
45,344,891	59,660	66	13,589,490	66,121	18,775	12,795,281	59,451	86	src/sqlite3.c:sqlite3VdbeRecordUnpack	
36,550,248	47,192	94	9,615,816	217,845	11,567	0	0	0	src/sqlite3.c:cellSizePtr	
35,156,491	1,031,905	859	7,810,853	489,509	1,936	6,546,085	175,469	26,159	/build/glibc-2.19/malloc/malloc.c:_int_malloc	
34,402,967	219,015	40	12,316,213	31,625	1,007	0	0	0	src/sqlite3.c:vdbeRecordCompareInt	
31,287,698	269,233	121	10,094,976	398,015	57,982	10,094,976	797,005	41,768	/build/glibc-2.19/string/../ports/sysdeps/aarch64/memcpy.S:memcpy	
30,895,222	1,055,479	718	3,990,072	45,246	157	3,247,672	1,200	58	src/sqlite3.c:sqlite3VXPrintf	
29,633,734	87	87	6,992,348	510,654	147,437	1,945,350	292	14	src/sqlite3.c:vdbeSorterSort	
28,301,654	1,222,726	236	7,685,792	129,350	101	4,693,862	15,480	91	src/sqlite3.c:sqlite3BtreeInsert	
27,452,670	605,975	428	7,719,336	275,711	3,045	6,130,240	1,247	180	/build/glibc-2.19/malloc/malloc.c:_int_free	
26,152,338	93,230	53	5,107,641	26,455	59	3,502,857	6,705	2	src/sqlite3.c:sqlite3VdbeSerialGet	
21,638,172	664,339	241	7,621,765	197,153	7,033	5,509,634	12,988	53	src/sqlite3.c:sqlite3PagerAcquire	
19,904,842	811,018	134	6,875,142	93,695	809	4,223,778	6,655	72	src/sqlite3.c:insertCell	
17,184,877	622,046	254	5,927,277	207,045	101	3,228,818	13,564	78	src/sqlite3.c:pager_write	
16,511,495	127,072	29	5,189,327	7,164	1,105	2,358,785	0	0	src/sqlite3.c:serialGet	
14,566,464	347,254	101	5,076,192	68,798	4,135	3,972,672	131,226	9,179	src/sqlite3.c:moveToChild	
14,089,915	528,334	433	3,522,612	169,118	295	1,089	24	22	???::???	
13,516,049	315,369	75	3,660,565	70,941	104	2,252,728	2,740	20	/build/glibc-2.19/malloc/malloc.c:malloc	
13,444,711	370,614	60	3,136,255	74,755	57,116	3,757,149	0	0	src/sqlite3.c:btreeParseCellIPtr	
11,814,468	620,489	364	3,444,231	109,318	159	1,401,768	11,253	7	src/sqlite3.c:sqlite3VdbeHalt	
9,867,819	655,851	130	3,350,976	68,237	46	1,820,276	62,050	70	src/sqlite3.c:moveToRoot	
9,023,249	615,625	175	2,774,458	27,649	72	1,719,012	578	1	src/sqlite3.c:sqlite3VdbeMemGrow	
9,015,155	136,420	114	2,528,161	33,460	40	1,808,361	12	7	/build/glibc-2.19/nptl/ptthread_mutex_lock.c:pthread_mutex_lock	
8,932,696	193,491	71	1,956,326	55,921	22	1,411,634	2	0	/build/glibc-2.19/nptl/ptthread_mutex_unlock.c:pthread_mutex_unlock	
8,916,165	82,925	47	2,092,310	0	0	1,933,573	1,583	3	src/sqlite3.c:memjournalWrite	
8,869,488	284,528	72	4,276,902	299,688	8,315	1,834,026	6,712	17	src/sqlite3.c:pcache1Fetch	
8,120,421	171,173	145	0	0	0	4,459,287	446,962	23,788	/build/glibc-2.19/string/../ports/sysdeps/aarch64/memset.S:memset	
7,759,659	338,888	58	2,364,882	24,321	1,308	1,624,112	104,416	1,631	src/sqlite3.c:sqlite3PcacheRelease	
6,799,934	97,805	282	2,068,211	38,793	684	1,555,697	3,672	11	src/sqlite3.c:sqlite3BtreeNext	
6,674,044	88,515	123	1,706,065	4,244	43	1,094,451	7	0	src/sqlite3.c:freeSpace	
6,536,765	760,083	320	2,119,849	121,314	85	1,091,200	0	0	src/sqlite3.c:sqlite3_step	

Performance Analysis with Linux Perf

Two modes of operation:

- ▶ sum up all counters: `perf stat`
- ▶ record events: `perf record`

Linux Perf: Example – SQLite

```
$ perf stat ./sqlite_llvm <test.sql >/dev/null
```

```
Performance counter stats for './sqlite_llvm':
```

1045.856070	task-clock (msec)	#	1.000 CPUs utilized	
1	context-switches	#	0.001 K/sec	
0	cpu-migrations	#	0.000 K/sec	
809	page-faults	#	0.774 K/sec	
1,636,720,010	cycles	#	1.565 GHz	[83.16%]
548,530,227	stalled-cycles-frontend	#	33.51% frontend cycles idle	[83.16%]
218,991,051	stalled-cycles-backend	#	13.38% backend cycles idle	[67.04%]
3,385,841,295	instructions	#	2.07 insns per cycle	
		#	0.16 stalled cycles per insn	[83.54%]
709,436,490	branches	#	678.331 M/sec	[83.54%]
2,586,354	branch-misses	#	0.36% of all branches	[83.17%]

```
1.045918998 seconds time elapsed
```

Linux Perf: Example – 483.xalancbmk

```
$ perf record ./xalancbmk
```

```
$ perf report
```

```
0.20 629a84:    ldr    w9, [x0,#24]
18.71 629a88:    ldr    w8, [x1,#24]
12.93 629a8c:    cmp     w9, w8
2.74 629a90:    b.ne   629af8 <xalanc_1_8::XalanDOMString::equals
2.00 629a94:    ldp     x8, x10, [x0]
2.43 629a98:    cmp     x8, x10
1.80 629a9c:    ldp     x10, x12, [x1]
1.03 629aa0:    adrp    x11, 704000 <vtable for xalanc_1_8::ReusableArenaBlock+0x8>
0.53 629aa4:    add     x11, x11, #0xb08
0.03 629aa8:    csel    x8, x11, x8, eq
1.33 629aac:    cmp     x10, x12
0.34 629ab0:    csel    x10, x11, x10, eq
1.78 629ab4:    cbz     w9, 629b00 <xalanc_1_8::XalanDOMString::equals
0.02 629ab8:    ldrh    w11, [x8]
4.02 629abc:    ldrh    w12, [x10]
3.75 629ac0:    cmp     w11, w12
1.03 629ac4:    b.ne   629b08 <xalanc_1_8::XalanDOMString::equals
1.16 629ac8:    lsl     x9, x9, #1
        629acc:    add     x8, x8, #0x2
        629ad0:    add     x10, x10, #0x2
        629ad4:    sub     x9, x9, #0x2
10.18 629ad8:    cbz     x9, 629b00 <xalanc_1_8::XalanDOMString::equals
0.01 629adc:    ldrh    w11, [x8],#2
18.79 629ae0:    sub     x9, x9, #0x2
0.00 629ae4:    ldrh    w12, [x10],#2
9.22 629ae8:    cmp     w11, w12
5.11 629aec:    b.eq   629ad8 <xalanc_1_8::XalanDOMString::equals
        629af0:    mov     w0, wzr
        629af4:    ret
0.69 629af8:    mov     w0, wzr
0.09 629afc:    ret
        629b00:    orr     w0, wzr, #0x1
0.10 629b04:    ret
        629b08:    mov     w0, wzr
        629b0c:    ret
```