

# Performance Analysis and Optimization of C++ Standard Libraries

Sebastian Pop, Aditya Kumar

SARC: Samsung Austin R&D Center

October 26, 2017

# Agenda

- ▶ C++ standard template libraries
- ▶ software performance analysis
- ▶ improvements to libc++ and libstdc++ performance

# C++ Standard Template Libraries

- ▶ STL is easy to use
  - ▶ standard interface: portable
  - ▶ easy to change data types: list, vector, deque, map, etc.
  - ▶ easy to change algorithms: iterators
- ▶ complexity of operations specified by standard
- ▶ performance left to implementation

# Performance of STL implementations

- ▶ performance and memory usage depend on
  - ▶ implementation: libc++ vs. libstdc++ vs. MSVC, etc.
  - ▶ container type
  - ▶ inefficiencies in implementations
- ▶ always analyze software performance to validate your choice

Figure: Size in bytes of empty containers on x86\_64

Container	libstdc++	libc++	MSVC
vector<int>	24	24	24
list<int>	24	24	16
deque<int>	80	48	40
set<int>	48	24	16
unordered_set<int>	56	40	64
map<int, int>	48	24	16
unordered_map<int, int>	56	40	64

# Software Performance Analysis

- ▶ identify hot functions from execution profiles
- ▶ inspect hot path: unit-benchmarking
- ▶ identify resource utilization on hot path

## Profiling: identify hot path

- ▶ linux-perf, oprofile: cycles, instructions, HW counters
- ▶ valgrind: cachegrind (R/W/Instrs), callgrind

# Unit-benchmarking: inspect hot path

unit-benchmarking is unit-testing for performance

- ▶ set-up data structures in memory
- ▶ time hot function
- ▶ execute hot function until performance measures stabilize <sup>1</sup>

check performance of a single hot operation: less noise, keep focus

---

<sup>1</sup><https://github.com/google/benchmark>

## Example: performance analysis of std::string.find()

```
$ git clone https://github.com/hiraditya/std-benchmark.git && cd std-benchmark
$ git submodule update --recursive --remote
$ mkdir build && cd build && cmake .. && make -j8
$ cat ../cxx/string.bench.cpp
[...]  
    while (state.KeepRunning()) {  
        pos = s1.find(s2);  
        benchmark::DoNotOptimize(pos);  
    }  
[...]  
$ perf record ./cxx/string.bench.cpp.out --benchmark_filter=BM_find/16384  
$ perf report  
  
79.71% libstdc++.so.6.0.21    [.] std::__cxx11::basic_string<char, [...] >::find  
  
14.08 |78:    add    $0x1,%rbx  
13.82 |    add    $0x1,%rbp  
12.74 |    cmp    %r15,%rax  
      |    ja     b8  
18.21 |85:    cmp    -0x1(%rbp),%r14b  
13.16 |    lea    -0x1(%rbx),%r13  
13.28 |    mov    %rbx,%rax  
14.62 |    jne    78
```



# Analyze resource utilization on hot path

Inspect:

- ▶ source code, compiler IR, assembly code
- ▶ CPU usage, instructions used and their latencies
- ▶ memory bus and caches: loads/stores, spills, cache misses

# Improve Software Performance

- ▶ eliminate unnecessary work
  - ▶ call functionality from libc or libc++
  - ▶ reduce bus traffic: vectorize loads and stores
  - ▶ help compiler remove redundancies: attributes and inline
- ▶ analyze performance of different implementations
  - ▶ change data structures
  - ▶ change algorithms
  - ▶ change STL implementations
- ▶ analyze trade-offs of caching previous results
  - ▶ use more memory vs. less computation (and vice versa)

# Our contributions to libc++ and libstdc++

- ▶ `std::string.find()`: find string within string
- ▶ `xsgetn` in libc++: string to int value parsing
- ▶ inline several ctors and dtors
- ▶ add attribute `noreturn` to non-returning functions
- ▶ fixed quadratic behavior of `std::sort` of libc++

# Issues with std::string.find (libc++ and libstdc++)

b1, e1 iterators to the haystack string

b2, e2 iterators to the needle string

```
__search(b1, e1, b2, e2) {
```

```
...
```

```
while (true)
```

```
{
```

```
    while (true)
```

```
    {
```

```
        if (__first1 == __s)
```

```
            return make_pair(__last1, __last1);
```

```
        if (__pred(*__first1, *__first2))
```

```
            break;
```

```
        ++__first1;
```

```
    }
```

```
    _RandomAccessIterator1 __m1 = __first1;
```

```
    _RandomAccessIterator2 __m2 = __first2;
```

```
    while (true)
```

```
    {
```

```
        if (++__m2 == __last2)
```

```
            return make_pair(__first1, __first1 + __len2);
```

```
        ++__m1;
```

```
        if (!__pred(*__m1, *__m2))
```

```
        {
```

```
            ++__first1;
```

```
            break;
```

```
        }
```

```
    }
```

```
}
```

```
}
```

```
...
```

```
}
```

} Find the first matching character

} Match rest of the string

# Improved std::string.find (libc++ and libstdc++)

```
inline _LIBCPP_CONSTEXPR_AFTER_CXX11 const _CharT *
__search_substring(const _CharT *__first1, const _CharT *__last1, const _CharT *__first2, const _CharT *__last2) {
...
    // First element of __first2 is loop invariant.
    _CharT __f2 = *__first2;
    while (true) {
        __len1 = __last1 - __first1;
        // Check whether __first1 still has at least __len2 bytes.
        if (__len1 < __len2)
            return __last1;

        // Find __f2 the first byte matching in __first1.
        __first1 = _Traits::find(__first1, __len1 - __len2 + 1, __f2);
        if (__first1 == 0)
            return __last1;

        if (_Traits::compare(__first1, __first2, __len2) == 0)
            return __first1;

        ++__first1; // TODO: Boyer-Moore can be used.
    }
}
```

Find the first matching character

Match rest of the string

- ▶ replace byte by byte compare with call to memchr + memcmp
- ▶ 12x speedup on std-benchmark <sup>2</sup>

<sup>2</sup><https://github.com/hiraditya/std-benchmark>

# Our contributions to libc++

- ▶ string to int value parsing: `xsgetn` in libc++
  - ▶ replace byte by byte copy with call to libc `memcpy`
  - ▶ important speedup on proprietary benchmark
- ▶ inline ctor/dtor
  - ▶ `shared_ptr`
  - ▶ `basic_string`
- ▶ add attribute `noreturn` to non-returning functions
  - ▶ `__locale`, `vector`, `deque`, `future`, `regex`, `system_error`, etc.
  - ▶ important for compiler optimizations
  - ▶ remove false positives in static analysis tools

# Issues with `std::sort` (libc++)





Figure: Quadratic behavior of libc++ compared to libstdc++

\* [https://bugs.llvm.org/show\\_bug.cgi?id=20837](https://bugs.llvm.org/show_bug.cgi?id=20837)

## std::sort (libc++)

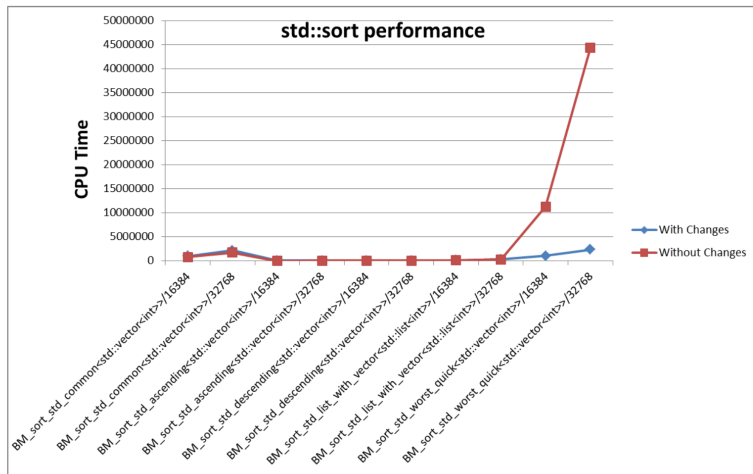
- ▶ Convert to introsort
- ▶ Sorting technique, which begins with quicksort and switches to heapsort after recursion reaches a threshold
- ▶ Worst case complexity of  $O(N\log N)$
- ▶ Eliminate recursion
- ▶ Replaced memory intensive recursive calls with stack  
std::stack uses std::deque, which uses std::algorithm
- ▶ Improved worst case time complexity by a factor of 10  
<https://reviews.llvm.org/D36423>
- ▶ quicksort with tail recursion elimination: quadratic worst case
- ▶ reimplemented as introsort: begin with quicksort, switch to heapsort when recursion depth goes beyond a threshold
- ▶ 16x speedup in the worst case (std-benchmark <sup>3</sup>)

---

<sup>3</sup><https://github.com/hiraditya/std-benchmark> 



# Sorting Results Plot (With std-benchmark)



# Questions